

Advanced image synthesis


Romain Vergne – 2014/2015



Ray tracing

Basically 2 functions:

- trace
 - find intersection with an object
- directIllumination
 - Direct lighting at a given point



Ray tracing

- `color trace(ray) {`
 - `hit = intersectScene(ray)`
 - `if(hit) {`
 - `color = directIllumination(hit)`
 - `if hit is reflective`
 - `color += c_refl * trace(reflected ray)`
 - `if hit is transmissive`
 - `color += c_trans * trace(refracted ray)`
 - `} else`
 - `color = background_color`
 - `return color`
- `}`



Ray tracing

- `color directIllumination(hit) {`
 - `color = (0,0,0)`
 - `for each light L {`
 - `T = cast shadow ray to L`
 - `if hit is not shadowed by L`
 - `color += Ambient+diffuse+specular terms(L, hit)`
 - `}`
 - `return color`
- `}`



Physics → Maths

- Bidirectionnal
- Reflectance
- Distribution
- Function

$$f(\mathbf{l}, \mathbf{v})$$



$$L_o(\mathbf{v}) = \int_{\Omega} f(\mathbf{l}, \mathbf{v}) \otimes L_i(\mathbf{l}) (\mathbf{n} \cdot \mathbf{l}) d\omega_i$$



Outgoing
radiance



BRDF



Ingoing
radiance



Surface
orientation


Reflectance equation



Microfacet theory


- Surface reflection (specular term)

$$f(\mathbf{l}, \mathbf{v}) = \frac{F(\mathbf{l}, \mathbf{h})G(\mathbf{l}, \mathbf{v}, \mathbf{h})D(\mathbf{h})}{4(\mathbf{n} \cdot \mathbf{l})(\mathbf{n} \cdot \mathbf{v})}$$





Microfacet theory

- Subsurface reflection (diffuse term)
 - Constant: $f_{\text{Lambert}}(\mathbf{l}, \mathbf{v}) = \frac{c_{\text{diff}}}{\pi}$



Textures


- From data
 - Colors, coefs, normals stored in 2D images
- Procedural shader
 - Little program that compute info at a given position




Textures



What is missing?

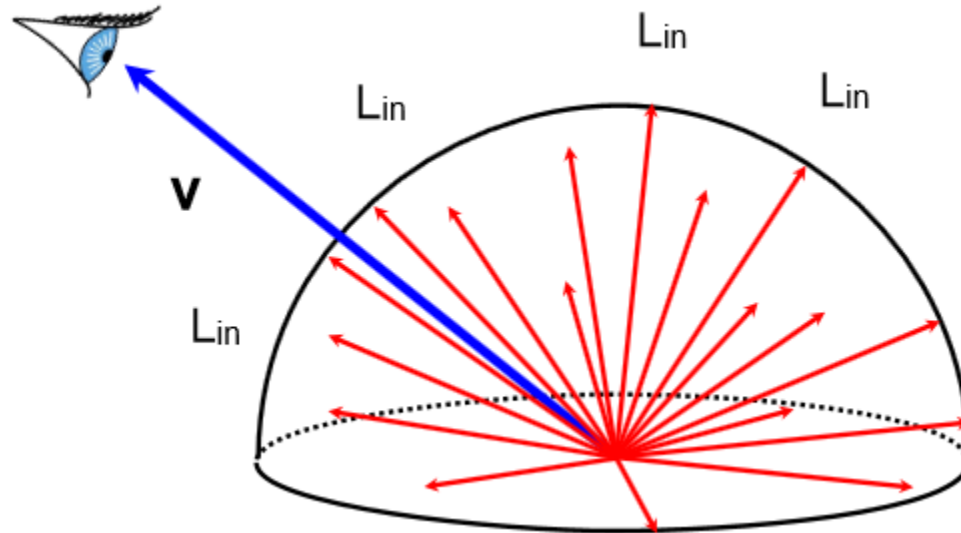


What is missing?



Reflectance equation

$$L_o(\mathbf{v}) = \int_{\Omega} f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l}) (\mathbf{n} \cdot \mathbf{l}) d\omega_i$$




- Where does L_{in} come from?



Reflectance equation

$$L_o(\mathbf{v}) = \int_{\Omega} f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l})(\mathbf{n} \cdot \mathbf{l}) d\omega_i$$




- Where does L_i come from?



Reflectance equation

$$L_o(\mathbf{v}) = \int_{\Omega} f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l})(\mathbf{n} \cdot \mathbf{l}) d\omega_i$$




- Where does L_i come from?



Rendering equation

$$L_o(\mathbf{v}) = \int_{\Omega} f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l})(\mathbf{n} \cdot \mathbf{l}) d\omega_i + E_{out}$$

[Kajiya 1986]



- Where does L_i come from?



Rendering equation


$$L_o(\mathbf{v}) = \int_{\Omega} f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l})(\mathbf{n} \cdot \mathbf{l}) d\omega_i + E_{out}$$

- Analytic solution usually impossible
- Lots of ways to solve it approximately:
 - e.g. ray-tracing – but approximation only...
 - How can we do better?




Soft contours (antialiasing)

- One ray per pixel
 - True or false...




Soft contours (antialiasing)

- One ray per pixel
 - True or false...
- Multiple rays per pixel
 - Average result (sum + divide by number of rays)




Soft contours (antialiasing)

No antialiasing




Soft contours (antialiasing)

16 samples / pixel





Soft shadows

- One sample per light
 - Hard (or noisy) shadows




Soft shadows

- One sample per light
 - Hard (or noisy) shadows
- Multiple samples per light
 - Soft shadows




Soft shadows

1 (random) sample / light




Soft shadows

1 (random) sample / light





Soft reflections (glossiness)

- One sample per hemisphere
 - Specular reflections




Soft reflections (glossiness)

- One sample per hemisphere
 - Specular reflections
- Multiple samples per hemisphere
 - Glossy materials



Soft reflections (glossiness)


- One sample per hemisphere
 - Specular reflections
- Multiple samples per hemisphere
 - Glossy materials




Soft reflections (glossiness)

- One sample per hemisphere
 - Specular reflections
- Multiple samples per hemisphere
 - Glossy materials

1 sample / hemisphere




50 sample / hemisphere



Soft refractions (translucency)

- One sample per hemisphere
 - Specular refractions
- Multiple samples per hemisphere
 - Translucent materials




Soft wavelengths (diffraction)

- Sampling wavelengths




Soft depth (focal length)

- Sampling aperture



Soft depth (focal length)

- Sampling aperture





© 2005 Christoph Marquardt



Soft depth (focal length)

- Sampling aperture in practice
 - Sample a circular region
 - In the direction of the focal point
 - Intersect focal plane at the direction in the center of the disk



Soft depth (focal length)

- Sampling aperture in practice
 - Sample a circular region
 - In the direction of the focal point
 - Intersect focal plane at the direction in the center of the disk



Soft animation (motion blur)

- Sampling time
 - Objects in motion (between time 1 and 2)



Summary

Distributed ray tracing

- Sampling pixels → antialiasing
- Sampling light → soft shadows
- Sampling wavelengths → dispersion
- Sampling aperture → depth of field
- Sampling time → motion blur
- Sampling reflection function → blurred reflections
- Sampling refraction function → blurred refractions
- Sampling paths → interreflections

$$L_o(\mathbf{v}) = \int f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l}) (\mathbf{n} \cdot \mathbf{l}) d\omega_i$$



Summary

Distributed ray tracing

- Sampling pixels → antialiasing
- Sampling light → soft shadows
- Sampling wavelengths → dispersion
- Sampling aperture → depth of field
- Sampling time → motion blur
- Sampling reflection function → blurred reflections
- Sampling refraction function → blurred refractions
- Sampling paths → interreflections

$$L_o(\mathbf{v}) = \int f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l})(\mathbf{n} \cdot \mathbf{l}) d\omega_i$$





$$L_o(\mathbf{v}) = \int \int \int \int \int \int \int f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l})(\mathbf{n} \cdot \mathbf{l}) d\omega_i dx dy du dv dt d\lambda$$





Sampling strategy

- Regular sampling



Sampling strategy

- Uniform sampling



Sampling strategy

- Stratified sampling





●		●	●
●	●	●	●
●	●	●	●
●	●	●	●



Sampling strategy

- Regular sampling
- Uniform sampling
- Stratified sampling
- And many others
 - Poisson disk sampling
 - Low discrepancy sampling
 - ...



Ray tracing produces realistic images



Path notation

- Each path begins and ends with:
 - E – the eye
 - L – the light



Path notation

- Each path begins and ends with:
 - E – the eye
 - L – the light
- Each bounce involves interaction with a surface:
 - D – diffuse reflection
 - G – glossy reflection
 - S – specular reflection



Path notation

- Each path begins and ends with:
 - E – the eye
 - L – the light
- Each bounce involves interaction with a surface:
 - D – diffuse reflection
 - G – glossy reflection
 - S – specular reflection
- Ray-casting:
 - $E(D,G)L$
- Ray-tracing:
 - $E[S^*](D,G)L$
- All paths:
 - $E(D,G,S)^*L$



Path notation

- Each path begins and ends with:
 - E – the eye
 - L – the light
- Each bounce involves interaction with a surface:
 - D – diffuse reflection
 - G – glossy reflection
 - S – specular reflection
- Ray-casting:
 - $E(D,G)L$
- Ray-tracing:
 - $E[S^*](D,G)L$
- All paths:
 - $E(D,G,S)^*L$

Color bleeding?
Caustics?



Mathematical theory

- Rendering equation is complex
 - No analytical solution

$$L_o(\mathbf{v}) = \int f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l})(\mathbf{n} \cdot \mathbf{l}) d\omega_i$$


- → numerical scheme
 - Monte-carlo = powerfull tool to solve integrals that do not have analytical solutions (such as the rendering equation)



Mathematical theory

- Discretization
 - Trapezoid
 - Simpson's rule
 - Etc...

$$I = \int_a^b f(x) dx \quad \Rightarrow \quad I \approx \frac{(b-a)}{N} \sum_i^N f(x_i)$$



Mathematical theory


- Monte-carlo

- Random samples

- Average

- Don't keep track on spacing (hope it will be $1/N$ on average)

$$I = \int_a^b f(x) dx \quad \Rightarrow \quad I \approx \frac{(b-a)}{N} \sum_i^N f(x_i)$$



Mathematical theory

- Monte-carlo
 - Random samples
 - Average
 - Don't keep track on spacing (hope it will be $1/N$ on average)

$$I = \int_a^b f(x) dx \quad \Rightarrow \quad I \approx \frac{(b-a)}{N} \sum_i^N f(x_i)$$

- MC estimator:
$$F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{1/(b-a)}$$



Mathematical theory

- Monte-carlo $I = \int_a^b f(x)dx \Rightarrow I \approx \frac{(b-a)}{N} \sum_i^N f(x_i)$
 - Random samples
 - Average
 - Don't keep track on spacing (hope it will be $1/N$ on average)

- MC estimator: $F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{1/(b-a)}$

- Expected value: $E[F_N] = \int_a^b f(x)dx$

- Variance: $\sigma^2[F_N] = E[F_N^2] - (E[F_N])^2$



Mathematical theory

- MC estimator:
$$F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{1/(b-a)}$$
- Generalization:
$$F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{pdf(x_i)}$$



Mathematical theory

- MC estimator: $F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{1/(b-a)}$

- Generalization: $F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{\text{pdf}(x_i)}$

Depends on the sampling strategy



Mathematical theory

- MC estimator: $F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{1/(b-a)}$

- Generalization: $F_N = \frac{1}{N} \sum_{i=1}^N \frac{f(x_i)}{\text{pdf}(x_i)}$


Depends on the sampling strategy

- Convergence rate: $O(\sqrt{N})$
 - Divide error by 2 needs 4x samples




Example: compute π

- Integral of the function that is 1 inside the circle and 0 outside
 - Probability = $\pi/4$



Example: compute PI

- Integral of the function that is 1 inside the circle and 0 outside
 - Probability = $\text{PI}/4$
- Sample random positions
 - Count ratio $n = \text{\#inside} / \text{\#total}$
 - $\text{PI} \approx 4 * n$



Pros / cons

- Slow for 1D problems, but:
 - Convergence is independent of dimensions
 - Few restrictions on the integrand
 - Conceptually simple
 - Efficient to solve at just a few points




Pros / cons

- Slow for 1D problems, but:
 - Convergence is independent of dimensions
 - Few restrictions on the integrand
 - Conceptually simple
 - Efficient to solve at just a few points
- Cons:
 - Noisy
 - Slow convergence
 - Good implementation is hard




Monte-carlo ray-tracing

- One ray per pixel




Monte-carlo ray-tracing

- One ray per pixel
- On each point, generate random rays, accumulate radiance




Monte-carlo ray-tracing

- One ray per pixel
- On each point, generate random rays, accumulate radiance
- Recurse




Monte-carlo ray-tracing

- One ray per pixel
- On each point, generate random rays, accumulate radiance
- Recurse
- Systematically sample light




Monte-carlo ray-tracing




Monte-carlo path-tracing

- One single ray per bounce




Monte-carlo path-tracing

- One single ray per bounce
- But hundreds of ray per pixel




Monte-carlo path-tracing

- 10 paths / pixel




Monte-carlo path-tracing

- 10 paths / pixel



Monte-carlo path-tracing

- 100 paths / pixel



Importance sampling



Naïve sampling strategy




Optimal sampling strategy



Importance sampling

- Sampling strategy
 - Non-uniform sampling
 - Sample more in places where there are likely to be larger contributions to the integral



Importance sampling

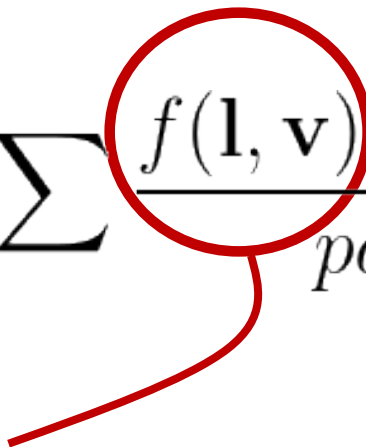
- Sampling strategy
 - Non-uniform sampling
 - Sample more in places where there are likely to be larger contributions to the integral

$$L_o(\mathbf{v}) \approx \frac{1}{N} \sum \frac{f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l})(\mathbf{n} \cdot \mathbf{l})}{pdf(\mathbf{l})}$$



Importance sampling

- Sampling strategy
 - Non-uniform sampling
 - Sample more in places where there are likely to be larger contributions to the integral

$$L_o(\mathbf{v}) \approx \frac{1}{N} \sum \frac{f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l}) (\mathbf{n} \cdot \mathbf{l})}{pdf(\mathbf{l})}$$


Send more rays in the direction of reflection
(depending on glossiness properties)



Importance sampling

- Sampling strategy
 - Non-uniform sampling
 - Sample more in places where there are likely to be larger contributions to the integral

$$L_o(\mathbf{v}) \approx \frac{1}{N} \sum \frac{f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l}) (\mathbf{n} \cdot \mathbf{l})}{pdf(\mathbf{l})}$$

Avoid rays at grazing angles

Send more rays in the direction of reflection
(depending on glossiness properties)



Importance sampling

- Sampling strategy
 - Non-uniform sampling
 - Sample more in places where there are likely to be larger contributions to the integral

$$L_o(\mathbf{v}) \approx \frac{1}{N} \sum \frac{f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l}) (\mathbf{n} \cdot \mathbf{l})}{pdf(\mathbf{l})}$$

Avoid rays at grazing angles

Send more rays in the direction of reflection
(depending on glossiness properties)



Importance sampling

- Sampling strategy
 - Non-uniform sampling
 - Sample more in places where there are likely to be larger contributions to the integral

$$L_o(\mathbf{v}) \approx \frac{1}{N} \sum \frac{f(\mathbf{l}, \mathbf{v}) L_i(\mathbf{l}) (\mathbf{n} \cdot \mathbf{l})}{pdf(\mathbf{l})}$$

Send more rays in the direction of reflection
(depending on glossiness properties)

Avoid rays at grazing angles

Do not forget pdf



Uniform sampling the hemisphere = $1/\text{vol}(S) = 1/2\pi$



Importance sampling

- Sampling the hemisphere

$$(x \in [0, 1), y \in [0, 1)) \rightarrow \left(\phi = 2\pi x, \theta = \cos^{-1} \left[(1 - y)^{1/(m+1)} \right] \right)$$




Loi du cosinus surélevé

$$\mathbf{p} = \sin \theta \cos \phi \mathbf{u} + \sin \theta \sin \phi \mathbf{v} + \cos \theta \mathbf{w}$$




Importance sampling


- Sampling the hemisphere




$m = 1000$




$m = 100$




$m = 10$



$m = 1$



$m = 1$




$m = 0$



Importance sampling


- Sampling the BRDF




Importance sampling

- Sampling the BRDF


$$U(\omega_i)$$



5 Samples/Pixel



$$P(\omega_i)$$




Path-tracing extensions

- Bidirectional path tracing [Lafortune, Veach]
 - Combine lights from light and camera
- Metropolis [Veach]
 - Extension to bi-directional
 - Reuse and mutation of important paths




Caching approaches

- Indirect lighting is mostly smooth




Irradiance caching

- Indirect lighting is mostly smooth




Irradiance caching

- Indirect lighting is mostly smooth




Irradiance caching

- Indirect lighting is mostly smooth
- Interpolate nearby values





Irradiance caching

- Indirect lighting is mostly smooth
- Interpolate nearby values
- But compute full direct lighting



Irradiance caching

- Indirect lighting is mostly smooth
- Interpolate nearby values
- But compute full direct lighting




Irradiance caching

- Introduced by Ward [92]
- Lots of extensions
 - Gradient [92], SH [05], Hessian [12],...
- Complementary to photon mapping




Photon mapping

- Preprocess: cast rays from light sources, and bounce randomly
- Store photons: position, power, direction




Photon mapping

- Preprocess: cast rays from light sources, and bounce randomly
- Store photons: position, power, direction




Photon mapping

- Efficiently store photons for fast access
- Use spatial hierarchical data structure (kd-tree)




Photon mapping (rendering)

- Cast primary rays
- Secondary rays
 - Reconstruct irradiance using adjacent stored photons (k-closest photons)
- Combine with irradiance caching or other technique...




Photon mapping (rendering)

- Cast primary rays
- Secondary rays
 - Reconstruct irradiance using adjacent stored photons (k-closest photons)
- Combine with irradiance caching or other technique...





Photon mapping (rendering)

- Special photon map for caustics LS+D



Radiosity

- Diffuse hypothesis
- Discretize scene in patches



Radiosity

- Convert rendering equation in the form:

$$B_i = B_{ei} + f_i \sum_j F_{ij} B_j$$

- B_i = radiosity of patch i

- F_{ij} = form factor :
$$F_{ij} = \int_{A_i} \int_{A_j} v(\mathbf{x}, \mathbf{x}') \frac{\cos(\theta) \cos(\theta')}{\pi r^2} d\mathbf{x} d\mathbf{x}'$$



Radiosity

- Matricial equation

$$\begin{bmatrix} B_1 \\ \vdots \\ B_n \end{bmatrix} = \begin{bmatrix} B_{e,1} \\ \vdots \\ B_{e,n} \end{bmatrix} + \begin{bmatrix} f_1 F_{1,1} & \dots & f_1 F_{1,n} \\ \vdots & \ddots & \vdots \\ f_n F_{n,1} & \dots & f_n F_{n,n} \end{bmatrix} \begin{bmatrix} B_1 \\ \vdots \\ B_n \end{bmatrix}$$



Radiosity

- Matricial equation

$$\begin{bmatrix} B_1 \\ \vdots \\ B_n \end{bmatrix} = \begin{bmatrix} B_{e,1} \\ \vdots \\ B_{e,n} \end{bmatrix} + \begin{bmatrix} f_1 F_{1,1} & \dots & f_1 F_{1,n} \\ \vdots & \ddots & \vdots \\ f_n F_{n,1} & \dots & f_n F_{n,n} \end{bmatrix} \begin{bmatrix} B_1 \\ \vdots \\ B_n \end{bmatrix}$$

- - Long / memory / discontinuities (geom) / diffuse only
- + view independant / handle complex scenes /
- Mostly used in architecture and video games (light maps)



Radiosity

- More recent methods
 - Precomputed radiance transfer
 - Many lights (VPLs) and extensions



References

- MIT:
 - <http://ocw.mit.edu/courses/electrical-engineering-and-computer-science/6-837-computer-graphics-fall-2012/lecture-notes/>
- Stanford:
 - <http://candela.stanford.edu/cs348b-14/doku.php>
- Siggraph:
 - <http://blog.selfshadow.com/publications/s2014-shading-course/>
 - <http://blog.selfshadow.com/publications/s2013-shading-course/>
- Image synthesis & OpenGL:
 - http://romain.vergne.free.fr/blog/?page_id=97
- Path tracing and global illum:
 - <http://www.graphics.stanford.edu/courses/cs348b-01/course29.hanrahan.pdf>
 - http://web.cs.wpi.edu/~emmanuel/courses/cs563/write_ups/zackw/realistic_raytracing.html
- GLSL / Shadertoy:
 - <https://www.opengl.org/documentation/glsl/>
 - <https://www.shadertoy.com/>
 - <http://www.iquilezles.org/>
- <http://fileadmin.cs.lth.se/cs/Education/EDAN30/lectures/L2-rt.pdf>
- <http://csokavar.hu/raytrace/imm6392.pdf>
- http://web.cs.wpi.edu/~emmanuel/courses/cs563/write_ups/zackw/realistic_raytracing.html

